

CS 229, Summer 2025

Problem Set #1

Baran Peters (baranp@stanford.edu)

Due Friday, July 11, 2025 at 11:59 pm on Gradescope.

Notes: (1) These questions require thought, but do not require long answers. Please be as concise as possible.

(2) If you have a question about this homework, we encourage you to post your question on our Ed forum, at <https://edstem.org/us/courses/79825/discussion>.

(3) If you missed the first lecture or are unfamiliar with the collaboration or honor code policy, please read the policy on the course website before starting work.

(4) For the coding problems, you may not use any libraries except those defined in the provided `environment.yml` file. In particular, ML-specific libraries such as scikit-learn are not permitted.

(5) The due date is Friday, July 11, 2025 at 11:59 pm. If you submit after Friday, July 11, 2025 at 11:59 pm, you will begin consuming your late days. The late day policy can be found in the course website: Course Logistics and FAQ.

All students must submit an electronic PDF version of the written question including plots generated from the codes. We **require you to typeset your solutions via L^AT_EX**. All students must also submit a zip file of their source code to Gradescope, which should be created using the `make_zip.py` script. You should make sure to (1) restrict yourself to only using libraries included in the `environment.yml` file, and (2) make sure your code runs without errors. Your submission may be evaluated by the auto-grader using a private test set, or used for verifying the outputs reported in the writeup. Please make sure that your PDF file and zip file are submitted to the corresponding Gradescope assignments respectively. We reserve the right to not give any points to the written solutions if the associated code is not submitted.

Honor code: We strongly encourage students to form study groups. Students may discuss and work on homework problems in groups. However, each student must write down the solutions independently, and without referring to written notes from the joint session. In other words, each student must understand the solution well enough in order to reconstruct it by themselves. In addition, each student should write on the problem set the set of people with whom s/he collaborated. Further, because we occasionally reuse problem set questions from previous years, we expect students not to copy, refer to, or look at the solutions in preparing their answers. It is an honor code violation to intentionally refer to a previous year's solutions.

Regarding Notation: The notation used in this problem set matches the notation used in the lecture notes. Some notable differences from lecture notation:

- The superscript “ (i) ” represents an index into the training set – for example, $x^{(i)}$ is the i -th feature vector and $y^{(i)}$ is the i -th output variable. In lecture notation, these would instead be expressed as x_i and y_i .
- The subscript j represents an index in a vector – in particular, $x_j^{(i)}$ represents the j -th feature in the feature vector $x^{(i)}$. In lecture notation, $x_j^{(i)}$ would be $h_j(x_i)$ or $x_i[j]$.
- The vector that contains the weights parameterizing a linear regression is expressed by the variable θ , whereas lectures use the variable $\mathbf{w}$. As such, $\theta_0 = w_0$, $\theta_1 = w_1$, and so on.

An overview of this notation is also given at the beginning of the lecture notes.

1. [30 points] Convergence and Learning Rate for Gradient Descent

When running an optimization algorithm, finding a good learning rate may require some tuning. If the learning rate is too high, the gradient descent (GD) algorithm might not converge. If the learning rate is too low, GD may take too long to converge. In this question, we will investigate some theoretical guarantees for the convergence of GD on convex objective functions, and their implications on choosing the learning rate.

Recap and notation. Suppose we have a convex objective function $J(\theta)$. At each iteration, GD updates the iterate $\theta^{[t]}$ as follows:

$$\theta^{[t]} = \theta^{[t-1]} - \alpha \nabla J(\theta^{[t-1]})$$

where α is the learning rate and $\nabla J(\theta^{[t-1]})$ is the gradient of $J(\theta)$ evaluated at $\theta^{[t-1]}$.

(a) [8 points] Quadratic Objective, Scalar Variable

In this part, consider the simple objective function with one-dimensional variable θ :

$$J(\theta) = \frac{1}{2}\beta\theta^2$$

where $\theta \in \mathbb{R}$ is our parameter and $\beta > 0$ is a positive, constant scalar. We initialize GD at some initialization $\theta^{[0]} \neq 0$ and run GD with learning rate α .

- i. **Derive** the range of α such that the iterate of GD converges in the sense that there exists a scalar $\theta^\dagger$ such that $\lim_{t \rightarrow \infty} |\theta^{[t]} - \theta^\dagger| = 0$. Provide the value of $\theta^\dagger$ when GD converges and show that it's equal to the global minimum $\theta^* = \arg \min_{\theta} J(\theta)$. The range of α can potentially depend on β .
- ii. Given a desired accuracy $\epsilon > 0$, for all the α in the range that you derived, **derive** the minimum number of iterations T required to reach a point $\theta^{[T]}$ such that $|\theta^{[T]} - \theta^*| \leq \epsilon$. T can potentially depend on ϵ , $\theta^{[0]}$, α , and β . Investigate the behavior of T and whether T is increasing or decreasing for different values of α . Briefly discuss why choosing an inappropriate α can cause the algorithm to take longer to converge.

Hint: Express $\theta^{[t]}$ in terms of $\theta^{[0]}$, α , t , and β .

Hint: $\lim_{t \rightarrow \infty} c^t = 0, \quad \forall c \in \mathbb{R} \text{ s.t. } |c| < 1$

Answer:

i) To derive the range of α and determine $\Theta^{[t]}$, we need to first define the gradient of our objective function:

$$\nabla J(\Theta) = \beta\Theta$$

Therefore the GD Update function is as follows:

$$\Theta^{[t]} = \Theta^{[t-1]} - \alpha \cdot (\beta\Theta^{[t-1]})$$

We can also define $\Theta^{[t+1]}$ in terms of the series of updates starting from $\Theta^{[t]}$.

$$\Theta^{[t+1]} = (1 - \alpha\beta)\Theta^{[t]}$$

$$\Theta^{[t]} = (1 - \alpha\beta)\Theta^{[t-1]}$$

$$\Theta^{[1]} = (1 - \alpha\beta)\Theta^{[0]}$$

$$\Theta^{[2]} = (1 - \alpha\beta)\Theta^{[1]} = (1 - \alpha\beta)^2\Theta^{[0]}$$

$$\Theta^{[3]} = (1 - \alpha\beta)\Theta^{[2]} = (1 - \alpha\beta)^3\Theta^{[0]}$$

When continuing this series, we can generally say that:

$$\Theta^{[t]} = (1 - \alpha\beta)^t \Theta^{[0]}$$

We can now calculate the value of $\Theta^\dagger$ when Gradient Descent converges.

$$\lim_{t \rightarrow \infty} |\Theta^{[t]} - \Theta^\dagger| = 0$$

$$\lim_{t \rightarrow \infty} |(1 - \alpha\beta)^t \cdot \Theta^{[0]} - \Theta^\dagger| = 0$$

As $t \rightarrow \infty$, the term $(1 - \alpha\beta)^t \cdot \Theta^{[0]}$ converges towards zero.

$$\lim_{t \rightarrow \infty} \Theta^{[t]} = 0$$

$\Rightarrow$ So it follows $\Theta^\dagger = 0$.

Range of α so that GD converges.

For the above-mentioned limit to converge, the following needs to be true:

$$|1 - \alpha\beta| < 1$$

This is equivalent to:

$$-1 < 1 - \alpha\beta < 1$$

So:

$$\text{i) } -1 < 1 - \alpha\beta \iff \alpha\beta < 2 \iff \alpha < \frac{2}{\beta}$$

$$\text{ii) } 1 - \alpha\beta < 1 \iff -\alpha\beta < 0 \iff \alpha\beta > 0$$

It follows that for the range of α , we have $0 < \alpha < \frac{2}{\beta}$.

Global Minimum:

$$\nabla J(\Theta^*) = 0$$

$$\beta\Theta^* = 0 \quad | : \beta; \beta > 0$$

$$\Theta^* = 0$$

The global minimum Θ^* equals $\Theta^\dagger$.

ii)

Minimum Number of Iterations

For the minimum number of iterations T , the following is true:

$$|\Theta^{[T]} - \Theta^*| \leq \epsilon \iff |(1 - \alpha\beta)^T \cdot \Theta^{[0]} - \Theta^*| \leq \epsilon$$

With $\Theta^* = 0$, this simplifies to:

$$|(1 - \alpha\beta)^T \cdot \Theta^{[0]}| \leq \epsilon$$

With logarithms, this can also be expressed as:

$$\begin{aligned} |(1 - \alpha\beta)^T \cdot \Theta^{[0]}| &\leq \epsilon && \Big| : \Theta^{[0]}; \log(\dots) \\ |(1 - \alpha\beta)^T \cdot \Theta^{[0]}| &\leq \frac{\epsilon}{|\Theta^{[0]}|} \iff T \geq \frac{\log(\frac{\epsilon}{|\Theta^{[0]}|})}{\log(\frac{1}{|1 - \alpha\beta|})} \end{aligned}$$

The closer α is to 0 or $\frac{2}{\beta}$, the bigger the required number of iterations get. This is due to the logarithm in the denominator getting closer to zero, due to which T increases immensely in size, eventually reaching infinity. This corresponds also with the intuition that, that bigger step sizes overshoot the desired minimum, and take also a lot longer to converge as we hop from "side to side", and lower step sizes just not progressing towards the critical point.

(b) [4 points] **Quadratic Objective, d-dimensional Variable**

In this part of the question, consider the objective function with d -dimensional variable θ :

$$J(\theta) = \frac{1}{2} \sum_{i=1}^d \beta_i \theta_i^2$$

where θ_i 's are our parameters and $\beta_i \in \mathbb{R}$ s.t. $\beta_i > 0$ are positive, constant scalars. Assume that we start from some $\theta^{[0]}$ where $\theta_i^{[0]} \neq 0$ for all i and run GD with learning rate α . Derive the range of learning rate α such that GD converges in the sense that there exists a vector $\theta^\dagger \in \mathbb{R}^d$ such that $\lim_{t \rightarrow \infty} \|\theta^{[t]} - \theta^\dagger\|_2 = 0$. Provide the value of $\theta^\dagger$ when GD converges. The range of α can potentially depend on β_i where $i \in \{1, \dots, d\}$.

Answer:

Since Θ is now a multidimensional vector, we can for now look at the i -th component of our parameters (as the solutions extend to the other elements as well).

Analogous to (a), we can define $\Theta_i^{[t]}$ in terms of $\Theta_i^{[0]}$:

$$\Theta_i^{[t]} = (1 - \alpha\beta_i)^t \cdot \Theta_i^{[0]}.$$

Value of $\Theta^\dagger$:

$$\lim_{t \rightarrow \infty} |\Theta_i^{[t]} - \Theta_i^\dagger| = 0$$

$$\lim_{t \rightarrow \infty} |(1 - \alpha\beta_i)^t \cdot \Theta_i^{[0]} - \Theta_i^\dagger| = 0$$

This has the same shape as the limit we solved for the single dimension case. Hence, we can say that for all $i \in \{1, \dots, d\}$ the following holds true:

$$\lim_{t \rightarrow \infty} \Theta^{[t]} = 0$$

$$\Rightarrow \Theta^\dagger = 0.$$

Range of α :

Also like in a), for GD to converge, the following needs to be true:

$$|1 - \alpha\beta_i| < 1.$$

This leads to:

$$0 < \alpha < \frac{2}{\beta_i} \quad \text{for all } i \in \{1, \dots, d\}.$$

(c) [4 points] Coding Question: Quadratic Multivariate Objective

Let the objective function be

$$J(\theta) = \theta^\top A \theta$$

where $A \in \mathbb{R}^{2 \times 2}$ is a 2×2 real, positive semi-definite matrix. Do the following exercises:

- i. Implement the `update_theta` and `gradient_descend` function in `src/gd_convergence/experiment.py`. You can stop the GD algorithm when either of the following condition is satisfied:

- A. $|J(\theta^{[t]}) - J(\theta^{[t-1]})| < \epsilon$ where $\epsilon = 10^{-50}$ is given as the function parameter. This is when we assume the algorithm converged.
- B. $J(\theta^{[t]}) > 10^{20}$. This is to prevent an infinite loop when the algorithm does not converge.

To test your implementation, run `python src/gd_convergence/experiment.py`, which checks that your θ (approximately) converges to the optimal value.

Note that we have provided you a matrix $A = \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$ and $\theta^{[0]} = [-1, 0.5]$ at the beginning of the file for you to experiment with. Therefore, the objective function is a special case of the objective function in part (b) with dimension $d = 2$. Check if your theoretical derivation of the feasible range of the learning rate indeed matches empirical observations.

- ii. Now, suppose we rotate the matrix A , what do we observe? Plot the trajectories of the GD algorithm using the following learning rates: 0.05, 0.1, 0.2, 0.3, 0.4, 0.45, 0.5, and 1. We have provided the rotation and plotting function for you. You need to simply run `python src/gd_convergence/plotting.py lr1 lr2 lr3 ...` with `lr*` replaced with desired learning rates. Include the output file `trajectories.png` and `trajectories_rotated.png`, as well as a brief discussion on the printed output of `plotting.py` in the write-up.

Remark: Setting the learning rate too high will cause the objective function to not converge. The convergence properties of these learning rates are also rotational invariant. We will show this formally in the next part of the problem.

Answer:




Figure 1: Gradient descent trajectories with regular A .




Figure 2: Gradient descent trajectories with rotated A .

As predicted by our result in a2), we can see that with very small learning rates GD takes a lot of steps to converge, same as with larger learning rates after a specific point ($\alpha \gtrsim 0.3$), as they tend to overshoot. After $\alpha \gtrsim 0.5$ GD doesn't seem to converge anymore and either the loss stays constant or increases until we stop after hitting our thresholds

(d) [12 points] **Convergence of Gradient Descent for a General Convex Objective**

Now let's consider any convex, twice continuously differentiable objective function $J(\theta)$ that is bounded from below. Suppose that the largest eigenvalue of the Hessian matrix $H = \nabla^2 J(\theta)$ is less than or equal to $\beta_{\max}$ for all points θ .

Show that when running gradient descent, choosing a step size $0 < \alpha < \frac{1}{\beta_{\max}}$ guarantees that $J(\theta^{[t]})$ will converge to some finite value as t approaches infinity.¹ Specifically, you should derive an inequality for $J(\theta^{[t]})$ in terms of $J(\theta^{[t-1]})$, $\nabla J(\theta^{[t-1]})$, α , and $\beta_{\max}$, and show that the objective function is strictly decreasing during each iteration, i.e. $J(\theta^{[t]}) < J(\theta^{[t-1]})$.

Hint: You can assume that, by Taylor's theorem, the following statement is true:

$$J(y) = J(x) + \nabla J(x)^\top (y - x) + \frac{1}{2}(y - x)^\top \nabla^2 J(x + c(y - x))(y - x) \text{ for some } 0 \leq c \leq 1$$

Hint: The Hessian of a convex function is symmetric and positive semi-definite at any point θ , which means all of its eigenvalues are real and non-negative.

Hint: The Hessian matrix is symmetric. Consider using the spectral theorem introduced in problem 3 in homework 0.

Optional (no credit): Using the gradient descent inequality that you derived, show that the GD algorithm converges in the sense that $\lim_{t \rightarrow \infty} \|\nabla J(\theta^{[t]})\|_2^2 = 0$.²

Remark: This question suggests that a smaller learning rate should be applied if we believe the curvature of the objective function is big.

Answer:

With the Taylor Theorem, we can define $J(\theta^{[t]})$ like the following:

Let's assume $x = \theta^{[t-1]}$ and $y = \theta^{[t]}$ (the point we want to define out going from the information at point x)

$$J(y) = J(x) + \nabla J(x)^\top (y - x) + \frac{1}{2}(y - x)^\top \nabla^2 J(x + c(y - x))(y - x) \quad (1)$$

(for some $0 \leq c \leq 1$.)

We can substitute for $y - x = \theta^{[t]} - \theta^{[t-1]}$ based on the Gradient Descent Update Rule:

$$\theta^{[t]} - \theta^{[t-1]} = -\alpha \nabla J(\theta^{[t-1]}) \quad (2)$$

¹This definition of convergence is different from the definition in previous questions where we explicitly prove that the parameter θ converges to an optimal parameter θ^* . In this case, we do not know the optimal value θ^* , and it would be difficult to prove convergence using the previous definition.

²Note that $\lim_{t \rightarrow \infty} \|\nabla J(\theta^{[t]})\|_2^2 = 0$ does not necessarily mean that there exists a vector $\hat{\theta}$ such that $\theta^{[t]} \rightarrow \hat{\theta}$. (Even an 1-dimensional counterexample exists.) However, for convex functions, when $\theta^{[t]}$ stays in a bounded set, $\lim_{t \rightarrow \infty} \|\nabla J(\theta^{[t]})\|_2^2 = 0$ does imply that $J(\theta^{[t]}) \rightarrow \inf_{\theta} J(\theta)$ as $t \rightarrow \infty$. Therefore, $\lim_{t \rightarrow \infty} \|\nabla J(\theta^{[t]})\|_2^2 = 0$ is typically considered a reasonable definition for an optimization algorithm's convergence.

Substituting:

$$\begin{aligned}
 J(\theta^{[t]}) &= J(\theta^{[t-1]}) + \nabla J(\theta^{[t-1]})^T (-\alpha \nabla J(\theta^{[t-1]})) \\
 &\quad + \frac{1}{2} (-\alpha \nabla J(\theta^{[t-1]}))^T \nabla^2 J(x + c(y - x)) (-\alpha \nabla J(\theta^{[t-1]})) \\
 &= J(\theta^{[t-1]}) - \alpha \|\nabla J(\theta^{[t-1]})\|_2^2 \\
 &\quad + \frac{\alpha^2}{2} \nabla J(\theta^{[t-1]})^T \nabla^2 J(x + c(y - x)) \nabla J(\theta^{[t-1]})
 \end{aligned} \tag{3}$$

To make sure that the objective function is strictly decreasing, we need to prove that the positive term (Hessian term) is smaller than the negative term (scaled L2 norm of the gradient).

For this we can use the Spectral Theorem together with the fact that we know that the largest eigenvalue of $H = \nabla^2 J(x)$:

$$\text{Let } H = \nabla^2 J(x + c(y - x)) = U \Lambda U^T.$$

$$\text{Let } Z = \nabla J(\theta^{[t-1]}).$$

$$Z^T U \Lambda U^T Z = \nabla J(\theta^{[t-1]})^T \nabla^2 J(x + c(y - x)) \nabla J(\theta^{[t-1]})$$

$$\text{Let } W = U^T Z.$$

$$\Rightarrow W^T \Lambda W = \sum_{i=1}^d \lambda_i W_i^2$$

With $\lambda_i \leq \beta_{max}$, it follows that:

$$\sum_{i=1}^d \lambda_i W_i^2 \leq \beta_{max} \sum_{i=1}^d W_i^2 = \beta_{max} \|W\|^2$$

As U is an orthogonal matrix and does not change length during transformation, we can also take the length of Z instead of W :

$$Z^T U \Lambda U^T Z \leq \beta_{max} \|Z\|^2$$

Inserting this into (3), we get:

$$\begin{aligned}
 J(\theta^{[t]}) &\leq J(\theta^{[t-1]}) - \alpha \|\nabla J(\theta^{[t-1]})\|_2^2 + \frac{\alpha^2}{2} \beta_{max} \|\nabla J(\theta^{[t-1]})\|_2^2 \\
 &\leq J(\theta^{[t-1]}) - \alpha \|\nabla J(\theta^{[t-1]})\|_2^2 \left(1 - \frac{\alpha \beta_{max}}{2}\right)
 \end{aligned} \tag{4}$$

For $J(\theta^{[t-1]}) > J(\theta^{[t]})$ to hold true the term in the parenthesis needs to be positive.

$$1 \geq \frac{\alpha \beta_{max}}{2} \tag{5}$$

With $0 < \alpha < 1/\beta_{max}$, α 's maximizing value is $1/\beta_{max}$.

$$1 \geq \frac{\frac{1}{\beta_{max}} \beta_{max}}{2}$$

$$1 \geq 0.5$$

$\rightarrow$ With $0 < \alpha < \frac{1}{\beta_{max}}$, we can guarantee that our objective function is decreasing with every step.

(e) [2 points] **Learning Rate for Linear Regression**

Consider using GD on the LMS objective introduced in lecture:

$$J(\theta) = \frac{1}{2} \|X\theta - y\|_2^2$$

where $X \in \mathbb{R}^{n \times d}$ is the design matrix of our data, $\theta \in \mathbb{R}^d$ is our parameter, and $\vec{y} \in \mathbb{R}^n$ is the response variable.

Let $\beta_{\max}$ be the largest eigenvalue of $X^\top X$. Prove that for all $\alpha \in (0, \frac{1}{\beta_{\max}})$, GD with learning rate α satisfies that $J(\theta^{[t]})$ converges as $t \rightarrow \infty$.

Hint: You can invoke the statements in the part (d) (including the optional question in part (d)) to solve this part (even if you were not able to prove part (d)).

Remark: The conclusion here suggests that for linear regression, roughly speaking, you should use smaller learning rates if the scale of your data X is big or if the data points are correlated with each other, both of which will enable a large top eigenvalue of $X^\top X$.

Remark: Even though in many cases the LMS objective can be solved exactly by solving the normal equation as shown in lecture, it is still useful to be able to guarantee convergence when using the Gradient Descent algorithm in situations where it is difficult to solve for the optimal θ^* directly (such as having large amount of data making inverting X very expensive).

Answer:

Same as in d), we can define $J(\theta^{(t)})$ with the Taylor Theorem:

Let $x = \theta^{(t-1)}$

Let $y = \theta^{(t)}$

$$\begin{aligned} J(\theta^{(t)}) &= J(\theta^{[t-1]}) - \alpha \|\nabla J(\theta^{[t-1]})\|_2^2 \\ &\quad + \frac{\alpha^2}{2} \nabla J(\theta^{[t-1]})^T \nabla^2 J(x + c(y - x)) \nabla J(\theta^{[t-1]}) \end{aligned} \tag{1}$$

Gradient & Hessian of $J(\theta)$

For $k = \{d_1, \dots, d_m\}$, we can isolate the k-th parameter of θ , θ_k :

$$\begin{aligned} \nabla_{\theta_k} J(\theta) &= \sum_{i=0}^N X_{ik} (X_i \theta - y_i) \\ &= X_k^T (X\theta - y) = X_k^T X\theta - X_k^T y \end{aligned}$$

Expanding this to the multidimensional case:

$$\nabla J(\theta) = X^T X\theta - X^T y$$

The Hessian resulting out of this is:

$$\nabla^2 J(\theta) = X^T X$$

Inserting this into (1) gives us:

$$J(\theta^{(t)}) = J(\theta^{(t-1)}) + \nabla J(\theta^{(t-1)})^T X^T X \nabla J(\theta^{(t-1)})$$

As left-multiplying a matrix by its own transpose leads to a symmetric matrix $Z = X^T X$, we can take over the same inequality as in (4) in d).

$$J(\theta^{(t)}) \leq J(\theta^{(t-1)}) - \alpha \cdot \|\nabla J(\theta^{(t-1)})\|_2^2 \cdot \left(1 - \frac{\alpha \beta_{max}}{2}\right) \quad (2)$$

Here again, d)'s statement (5) follows, which proves that for all $\alpha \in (0, 2/L_{max})$ with GD converges as $t \rightarrow \infty$.

2. [25 points] Locally weighted linear regression

- (a) [10 points] Consider a linear regression problem in which we want to “weight” different training examples differently. Specifically, suppose we want to minimize

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n w^{(i)} \left(\theta^T x^{(i)} - y^{(i)} \right)^2.$$

In class, we worked out what happens for the case where all the weights (the $w^{(i)}$'s) are the same. In this problem, we will generalize some of those ideas to the weighted setting. We will assume $w^{(i)} > 0$ for all i .

- i. [2 points] Show that $J(\theta)$ can also be written

$$J(\theta) = (X\theta - y)^T W (X\theta - y)$$

for an appropriate matrix W , and where X and y are as defined in class. Clearly specify the value of each element of the matrix W .

- ii. [4 points] If all the $w^{(i)}$'s equal 1, then we saw in class that the normal equation is

$$X^T X \theta = X^T y,$$

and that the value of θ that minimizes $J(\theta)$ is given by $(X^T X)^{-1} X^T y$. By finding the derivative $\nabla_{\theta} J(\theta)$ and setting that to zero, generalize the normal equation to this weighted setting, and give the new value of θ that minimizes $J(\theta)$ in closed form as a function of X , W and y .

- iii. [4 points] Suppose we have a dataset $\{(x^{(i)}, y^{(i)}); i = 1 \dots, n\}$ of n independent examples, but we model the $y^{(i)}$'s as drawn from conditional distributions with different levels of variance $(\sigma^{(i)})^2$. Specifically, assume the model

$$p(y^{(i)} | x^{(i)}; \theta) = \frac{1}{\sqrt{2\pi}\sigma^{(i)}} \exp \left(-\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2(\sigma^{(i)})^2} \right)$$

That is, each $y^{(i)}$ is drawn from a Gaussian distribution with mean $\theta^T x^{(i)}$ and variance $(\sigma^{(i)})^2$ (where the $\sigma^{(i)}$'s are fixed, known, constants). Show that finding the maximum likelihood estimate of θ reduces to solving a weighted linear regression problem. State clearly what the $w^{(i)}$'s are in terms of the $\sigma^{(i)}$'s.

In other words, this suggests that if we have prior knowledge on the noise levels (the variance of the label $y^{(i)}$ conditioned on $x^{(i)}$) of all the examples, then we should use weighted least squares with weights depending on the variances.

Answer:

i)

$$\text{Let } X = \begin{bmatrix} \dots (x^{(1)})^T \dots \\ \dots (x^{(2)})^T \dots \\ \vdots \\ \dots (x^{(n)})^T \dots \end{bmatrix}, \quad Y = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(n)} \end{bmatrix}$$

As our weights are just multiplied in the summation of $J(\theta)$, it must be a diagonal matrix, defined as:

$$W = \begin{bmatrix} w_1 & & 0 \\ & \ddots & \\ 0 & & w_n \end{bmatrix}$$

Let the error E be:

$$E = X\theta - Y = \begin{bmatrix} (x^{(1)})^T \theta - y^{(1)} \\ \vdots \\ (x^{(n)})^T \theta - y^{(n)} \end{bmatrix}$$

The cost function $J(\theta)$ is:

$$\begin{aligned} (1) \quad J(\theta) &= \sum_{i=1}^n w^{(i)} ((x^{(i)})^T \theta - y^{(i)})^2 = \sum_{i=1}^n w^{(i)} e_i^T e_i \\ &= \frac{1}{2} E^T W E = \frac{1}{2} (X\theta - Y)^T W (X\theta - Y) \end{aligned}$$

(The scaling factor of 0.5 can be omitted, as it doesn't change the landscape of the loss. So it does fit the equation in the problem statement.)

ii)

Derivative of weighted $J(\theta)$:

$$\begin{aligned} \nabla_{\theta} J(\theta) &= \nabla_{\theta} \frac{1}{2} ((X\theta - Y)^T W (X\theta - Y)) \\ &= \nabla_{\theta} \frac{1}{2} (\theta^T X^T W X \theta - \theta^T X^T W Y - Y^T W X \theta) \\ &= \nabla_{\theta} \frac{1}{2} (\theta^T X^T W X \theta) - 2 \nabla_{\theta} (\theta^T X^T W Y) \end{aligned}$$

With $\nabla_z z^T A z = 2Az$ for $z = \theta$ and $A = X^T W X$, we can simplify even further:

$$\begin{aligned} \nabla_{\theta} J(\theta) &= \frac{1}{2} (2X^T W X \theta - 2X^T W Y) \\ &= X^T W X \theta - X^T W Y \end{aligned}$$

Find minimum of $J(\theta)$:

$$\begin{aligned} \nabla J(\theta) &= 0 \\ 0 &= X^T W X \theta - X^T W Y \\ X^T W X \theta &= X^T W Y \\ \theta &= (X^T W X)^{-1} X^T W Y = \arg \min_{\theta} J(\theta) \end{aligned}$$

iii)

The Maximum Likelihood Estimate requires us to find, given the data we observed, the set of parameters θ that makes it the most likely. For this, we assume that our dataset underlying distribution follows a normal distribution with mean $\theta^T x^{(i)}$ and variance $\sigma^{(i)}$. For this, we can use the so-called likelihood function, giving us the likelihood of observing a specific set of data for that parameter θ .

As each sample is independent we can also say:

$$L(\theta) = p(y|X; \theta) = \prod_{i=1}^k p(y^{(i)}|x^{(i)}; \theta) = \prod_{i=1}^k \frac{1}{\sqrt{2\pi}\sigma^{(i)}} \exp\left(-\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2(\sigma^{(i)})^2}\right)$$

To make finding $\operatorname{argmax} L(\theta)$ easier, we can use the log-probability $l(\theta)$:

$$\begin{aligned} l(\theta) &= \log(L(\theta)) = \sum_{i=1}^n \log(p(y^{(i)}|x^{(i)}; \theta)) \\ &= \sum_{i=1}^n \left(\log\left(\frac{1}{\sqrt{2\pi}\sigma^{(i)}}\right) - \frac{(y^{(i)} - \theta^T x^{(i)})^2}{2(\sigma^{(i)})^2} \right) \end{aligned}$$

To maximize, we need to minimize the second (and only) term that depends on θ :

$$\begin{aligned} &\operatorname{argmin}_{\theta} \left(\sum_{i=1}^n \frac{1}{2(\sigma^{(i)})^2} (y^{(i)} - \theta^T x^{(i)})^2 \right) \\ &= \operatorname{argmin}_{\theta} \left(\frac{1}{2} \sum_{i=1}^n \frac{1}{(\sigma^{(i)})^2} (y^{(i)} - \theta^T x^{(i)})^2 \right) \end{aligned}$$

The above expression looks exactly like the weighted least squares loss from (i). The $w^{(i)}$ term is equivalent to $\frac{1}{(\sigma^{(i)})^2}$ in the MLE.

(b) [10 points] **Coding problem.**

We will now consider the following dataset (the formatting matches that of Datasets 1-4, except $x^{(i)}$ is 1-dimensional):

`src/lwr/{train,valid,test}.csv`

In `src/lwr/lwr.py`, implement locally weighted linear regression using the normal equations you derived in Part (a) and using

$$w^{(i)} = \exp\left(-\frac{\|x^{(i)} - x\|_2^2}{2\tau^2}\right).$$

This is a fairly standard choice for weights, where the weight $w^{(i)}$ depends on the particular point x at which we're trying to evaluate y : if $|x^{(i)} - x|$ is small, then $w^{(i)}$ is close to 1; if $|x^{(i)} - x|$ is large, then $w^{(i)}$ is close to 0. Here, τ is the bandwidth parameter, and it controls how quickly the weight of a training example falls off with distance of its $x^{(i)}$ from the query point x .

Train your model on the `train` split using $\tau = 0.5$, then run your model on the `valid` split and report the mean squared error (MSE). Finally plot your model's predictions on

the validation set (plot the training set with blue 'x' markers and the predictions on the validation set with a red 'o' markers). Does the model seem to be under- or overfitting?

Answer:




Figure 3: Locally weighted linear regression with $\tau = 0.5$ on the 'train' split.
(MSE: 0.330531268213752)

→ The model seems to be underfitting the curve, it's not flexible enough to capture the underlying oscillating pattern.

(c) [5 points] **Coding problem.**

We will now tune the hyperparameter τ . In `src/lwr/tau.py`, find the MSE value of your model on the validation set for each of the values of τ specified in the code. For each τ , plot your model's predictions on the validation set in the format described in part (b). Report the value of τ which achieves the lowest MSE on the `valid` split, and finally report the MSE on the `test` split using this τ -value.

Answer:

The lowest validation loss has been achieved by $\tau = 0.05$ with ≈ 0.0124 MSE loss.

On the test split, it achieves a ≈ 0.0169 MSE loss

3. [25 points] Linear regression: linear in what?

In this question, we will see how linear regression can be used to fit non-linear functions of the data using feature maps, as well as exploring some of its limitations.

(a) [5 points] Learning degree-3 polynomials of the input

Suppose we have a dataset $\{(x^{(i)}, y^{(i)})\}_{i=1}^n$ where $x^{(i)}, y^{(i)} \in \mathbb{R}$. We would like to fit a third degree polynomial $h_\theta(x) = \theta_3 x^3 + \theta_2 x^2 + \theta_1 x + \theta_0$ to the dataset. The key observation here is that the function $h_\theta(x)$ is still linear in the unknown parameter θ , even though it's not linear in the input x . This allows us to convert the problem into a linear regression problem as follows.

Let $\phi : \mathbb{R} \rightarrow \mathbb{R}^4$ be a function that transforms the original input x to a 4-dimensional vector defined as

$$\phi(x) = \begin{bmatrix} 1 \\ x \\ x^2 \\ x^3 \end{bmatrix} \in \mathbb{R}^4 \quad (3)$$

Let $\hat{x} \in \mathbb{R}^4$ be a shorthand for $\phi(x)$, and let $\hat{x}^{(i)} \triangleq \phi(x^{(i)})$ be the transformed input in the training dataset. We construct a new dataset $\{(\phi(x^{(i)}), y^{(i)})\}_{i=1}^n = \{(\hat{x}^{(i)}, y^{(i)})\}_{i=1}^n$ by replacing the original inputs $x^{(i)}$'s by $\hat{x}^{(i)}$'s. We see that fitting $h_\theta(x) = \theta_3 x^3 + \theta_2 x^2 + \theta_1 x + \theta_0$ to the old dataset is equivalent to fitting a linear function $h_\theta(\hat{x}) = \theta_3 \hat{x}_3 + \theta_2 \hat{x}_2 + \theta_1 \hat{x}_1 + \theta_0$ to the new dataset because

$$h_\theta(x) = \theta_3 x^3 + \theta_2 x^2 + \theta_1 x + \theta_0 = \theta_3 \phi(x)_3 + \theta_2 \phi(x)_2 + \theta_1 \phi(x)_1 + \theta_0 = \theta^T \hat{x} \quad (4)$$

In other words, we can use linear regression on the new dataset to find parameters $\theta_0, \dots, \theta_3$. Please write down 1) the objective function $J(\theta)$ of the linear regression problem on the new dataset $\{(\hat{x}^{(i)}, y^{(i)})\}_{i=1}^n$ and 2) the update rule of the batch gradient descent algorithm for linear regression on the dataset $\{(\hat{x}^{(i)}, y^{(i)})\}_{i=1}^n$.

Terminology: In machine learning, ϕ is often called the feature map which maps the original input x to a new set of variables. To distinguish between these two sets of variables, we will call x the input **attributes**, and call $\phi(x)$ the **features**. (Unfortunately, different authors use different terms to describe these two things. In this course, we will do our best to follow the above convention consistently.)

Answer:

1) $J(\theta)$ with linear Regression on the dataset $\{\hat{x}^{(i)}, y^{(i)}\}_{i=1}^n$:

$$\text{Let } \hat{x} \in \mathbb{R}^d = \phi(x), \hat{x}^{(i)} = \phi(x^{(i)}), \phi(x) = \begin{bmatrix} 1 \\ x \\ x^2 \\ x^3 \end{bmatrix}, \hat{X} = \begin{bmatrix} \phi(x^{(1)})^T \\ \vdots \\ \phi(x^{(n)})^T \end{bmatrix}$$

$$J(\theta) = \frac{1}{2} \|\hat{X}\theta - y\|_2^2 \quad (1)$$

$$= \frac{1}{2} \sum_{i=1}^n (\theta^T \hat{x}^{(i)} - y^{(i)})^2 \quad (2)$$

$$= \frac{1}{2} \sum_{i=1}^n (\theta^T \phi(x^{(i)}) - y^{(i)})^2 \quad (3)$$

2) Update rule of Batch GD on dataset $\{\hat{x}^{(i)}, y^{(i)}\}_{i=1}^n$:

(1) $\theta^{t+1} = \theta^t - \alpha \nabla J(\theta)$

Gradient for regular LMS: $\nabla_{\theta} J(\theta) = X^T(X\theta - y)$

→ for our case $\{\hat{x}^{(i)}, y^{(i)}\}_{i=1}^n$: $\nabla J(\theta) = \hat{X}^T(\hat{X}\theta - y)$

Inserting into the update rule:

(2)* $\theta^{t+1} = \theta^t - \alpha(\hat{X}^T(\hat{X}\theta - y))$

(b) [5 points] **Coding question: degree-3 polynomial regression**

For this sub-question question, we will use the dataset provided in the following files:

`src/featuremaps/{train,valid,test}.csv`

Each file contains two columns: x and y . In the terminology described in the introduction, x is the attribute (in this case one dimensional) and y is the output label.

Using the formulation of the previous sub-question, implement linear regression with **normal equations** using the feature map of degree-3 polynomials. Use the starter code provided in `src/featuremaps/featuremap.py` to implement the algorithm.

Create a scatter plot of the training data, and plot the learnt hypothesis as a smooth curve over it. Submit the plot in the writeup as the solution for this problem.

Remark: Suppose $\hat{X}$ is the design matrix of the transformed dataset. You may sometimes encounter a non-invertible matrix $\hat{X}^T \hat{X}$. For a numerically stable code implementation, always use `np.linalg.solve` to obtain the parameters directly, rather than explicitly calculating the inverse and then multiplying it with $\hat{X}^T y$.

Answer:




Figure 4: Polynomial regression with degree 3

(c) [5 points] **Coding question: degree- k polynomial regression**

Now we extend the idea above to degree- k polynomials by considering $\phi : \mathbb{R} \rightarrow \mathbb{R}^{k+1}$ to be

$$\phi(x) = \begin{bmatrix} 1 \\ x \\ x^2 \\ \vdots \\ x^k \end{bmatrix} \in \mathbb{R}^{k+1} \quad (4)$$

Follow the same procedure as the previous sub-question, and implement the algorithm with $k = 3, 5, 10, 20$. Create a similar plot as in the previous sub-question, and include the hypothesis curves for each value of k with a different color. Include a legend in the plot to indicate which color is for which value of k .

Submit the plot in the writeup as the solution for this sub-problem. Observe how the fitting of the training dataset changes as k increases. Briefly comment on your observations in the plot. **Answer:**




Figure 5: Polynomial regression with multiple degrees

The higher the polynomial degree gets, and the more features we get, the higher the variance gets and the more the curve overfits. It can model now more information, but with that also starts to learn the noise that comes with each sample and doesn't generalize as well anymore (high variance, low bias)

(d) [5 points] **Coding question: other feature maps**

You may have observed that it requires a relatively high degree k to fit the given training data, and this is because the dataset cannot be explained (i.e., approximated) very well

by low-degree polynomials. By visualizing the data, you may have realized that y can be approximated well by a sine wave. In fact, we generated the data by sampling from $y = \sin(x) + \xi$, where ξ is noise with Gaussian distribution. Please update the feature map ϕ to include a sine transformation as follows:

$$\phi(x) = \begin{bmatrix} 1 \\ x \\ x^2 \\ \vdots \\ x^k \\ \sin(x) \end{bmatrix} \in \mathbb{R}^{k+2} \quad (5)$$

With the updated feature map, train different models for values of $k = 0, 1, 2, 3, 5, 10, 20$, and plot the resulting hypothesis curves over the data as before.

Submit the plot as a solution to this sub-problem. Compare the fitted models with the previous sub-question, and briefly comment about noticeable differences in the fit with this feature map. **Answer:**

The fitted lines seem to fit much better straightaway from $k = 0$ onwards, due to the existence of the sine mapping of x . At $k=20$, we can even see overfitting with the model being more sensitive to specific points and not generalizing anymore as well, whereas the other lines basically seem on top of each other.

After checking the learned parameters for different k , it also seems like the other features are 10^1 to 10^3 times smaller and on an absolute level negligibly small.




Figure 6: Polynomial regression with sine features

(e) [5 points] **Overfitting with expressive models and small data**

For the rest of the problem, we will consider a small dataset (a random subset of the dataset you have been using so far) with much fewer examples, provided in the following file:

`src/featuremaps/small.csv`

We will be exploring what happens when the number of features start becoming bigger than the number of examples in the training set. Run your algorithm on this small dataset using the following feature map

$$\phi(x) = \begin{bmatrix} 1 \\ x \\ x^2 \\ \vdots \\ x^k \end{bmatrix} \in \mathbb{R}^{k+1} \quad (6)$$

with $k = 1, 2, 5, 10, 20$.

Create a plot of the various hypothesis curves (just like previous sub-questions). Observe how the fitting of the training dataset changes as k increases. Submit the plot in the writeup and comment on what you observe.

Remark: The phenomenon you observe where the models start to fit the training dataset very well, but suddenly “goes wild” is due to what is called *overfitting*. The intuition to have for now is that, when the amount of data you have is small relative to the expressive capacity of the family of possible models (that is, the hypothesis class, which, in this case, is the family of all degree k polynomials), it results in overfitting.

Loosely speaking, the set of hypothesis function is “very flexible” and can be easily forced to pass through all your data points especially in unnatural ways. In other words, the model explains the noises in the training dataset, which shouldn’t be explained in the first place. This hurts the predictive power of the model on test examples.

Answer: Starting with $k = 1$ the fitted curve looks very simple (high bias, low variance), and is just a line. The higher k gets the higher the polynomial of the resulting fitted curve is, and the better it can try to approximate the data points.

From $k > 5$ on (the point where also the features & degree of polynomial $>$ than the data points), we can see that the curves just get more complex without any real improvement, and seemingly irrelevant areas seem to get more complex. The curves start to overfit (high variance, low bias)




Figure 7: Polynomial regression with sine features

4. [35 points] Implicit Regularization

Recall that in the overparameterized regime (where the number of parameters is larger than the number of samples), typically there are infinitely many solutions that can fit the training dataset perfectly, and many of them cannot generalize well (that is, they have large validation errors). However, in many cases, the particular optimizer we use (e.g., GD, SGD with particular learning rates, batch sizes, noise, etc.) tends to find solutions that generalize well. This phenomenon is called implicit regularization effect (also known as algorithmic regularization or implicit bias).

In this problem, we will look at the implicit regularization effect on two toy examples in the overparameterized regime: linear regression and a quadratically parameterized model. For linear regression, we will show that gradient descent with zero initialization will always find the minimum norm solution (instead of an arbitrary solution that fits the training data), and in practice, the minimum norm solution tends to generalize well. For a quadratically parameterized model, we will show that initialization and batch size also affect generalization.

- (a) [3 points] Suppose we have a dataset $\{(x^{(i)}, y^{(i)}); i = 1, \dots, n\}$ where $x^{(i)} \in \mathbb{R}^d$ and $y^{(i)} \in \mathbb{R}$ for all $1 \leq i \leq n$. We assume the dataset is generated by a linear model without noise. That is, there is a vector $\beta^* \in \mathbb{R}^d$ such that $y^{(i)} = (\beta^*)^\top x^{(i)}$ for all $1 \leq i \leq n$. Let $X \in \mathbb{R}^{n \times d}$ be the matrix representing the inputs (i.e., the i -th row of X corresponds to $x^{(i)}$) and $\vec{y} \in \mathbb{R}^n$ the vector representing the labels (i.e., the i -th row of $\vec{y}$ corresponds to $y^{(i)}$):

$$X = \begin{bmatrix} - & x^{(1)} & - \\ - & x^{(2)} & - \\ \vdots & \vdots & \vdots \\ - & x^{(n)} & - \end{bmatrix}, \quad \vec{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(n)} \end{bmatrix}.$$

Then in matrix form, we can write $\vec{y} = X\beta^*$. We assume that the number of examples is less than the number of parameters (that is, $n < d$).

We use the least-squares cost function to train a linear model:

$$J(\beta) = \frac{1}{2n} \|X\beta - \vec{y}\|_2^2. \quad (7)$$

In this sub-question, we characterize the family of global minimizers to Eq. (7). We assume that $XX^\top \in \mathbb{R}^{n \times n}$ is an invertible matrix. **Prove that** β achieves zero cost in Eq. (7) if and only if

$$\beta = X^\top (XX^\top)^{-1} \vec{y} + \zeta \quad (8)$$

for some ζ in the subspace orthogonal to all the data (that is, for some ζ such that $\zeta^\top x^{(i)} = 0, \forall 1 \leq i \leq n$.)

Note that this implies that there is an infinite number of β 's such that Eq. (7) is minimized. We also note that $X^\top (XX^\top)^{-1}$ is the pseudo-inverse of X , but you don't necessarily need this fact for the proof.

Answer:

Given:

- i. $XX^\top \in \mathbb{R}^{n \times n}$ is an invertible matrix.
- ii. Cost function: $J(\beta) = \frac{1}{2n} \|X\beta - Y\|_2^2$
- iii. We are in the overparameterized regime: $n < d$
- iv. Linear model of dataset: $y^{(i)} = (\beta)^\top x^{(i)}$ or $Y = X\beta$

To prove this, we can first untangle the iff into the following:

1)

If β has the form in the problem statement, then the training loss is zero:

- i. $\beta = X^T(XX^T)^{-1}Y + \zeta$
- ii. Loss Function: $X\beta - Y = X(X^T(XX^T)^{-1}Y + \zeta) - Y$

Setting ii. to zero.

$$\begin{aligned} 0 &= X(X^T(XX^T)^{-1}Y + \zeta) - Y \\ &= XX^T(XX^T)^{-1}Y + X\zeta - Y \end{aligned}$$

As ζ is orthogonal to the whole rowspace of X , $X\zeta$ is zero
 $\implies Y = XX^T(XX^T)^{-1}Y \implies Y = Y$

2)

Assuming the training loss is zero, what form has β ?

$$\begin{aligned} 0 &= J(\beta) \\ 0 &= \frac{1}{n} \|X\beta - Y\|_2^2 \\ 0 &= X\beta - Y \\ Y &= X\beta \end{aligned}$$

As the matrix X has more columns than rows ($n < d$), the result of a left-multiplication with X has one part that is mapped to the column space of X (which is Y) but also another part $\zeta \in \mathbb{R}^d$ that is part of X 's nullspace and maps hence to zero:

$$\begin{aligned} Y + X\zeta &= X\beta \quad (\text{Left-multiplying by } X^T(XX^T)^{-1} = X^{-1}) \\ X^T(XX^T)^{-1}Y + X^T(XX^T)^{-1}X\zeta &= \beta \\ \implies \beta &= X^T(XX^T)^{-1}Y + \zeta \end{aligned}$$

- (b) [3 points] We still work with the setup of part (a). Among the infinitely many optimal solutions of Eq. (7), we consider the *minimum norm* solution. Let $\rho = X^\top(XX^\top)^{-1}\vec{y}$. In the setting of (a), **prove that** for any β such that $J(\beta) = 0$, $\|\rho\|_2 \leq \|\beta\|_2$. In other words, ρ is the minimum norm solution.

Hint: As an intermediate step, you can prove that for any β in the form of Eq. (8),

$$\|\beta\|_2^2 = \|\rho\|_2^2 + \|\zeta\|_2^2.$$

Answer:

- (1) $\rho = X^T(XX^T)^{-1}Y$
- (2) $\beta = \rho + \zeta$

Proof of $\|\rho\|_2 \leq \|\beta\|_2$ for $\forall \beta : J(\beta) = 0$:

$$(3) \|\beta\|_2^2 = \sum_i (\beta^i)^2$$

Let's define $\beta^{(i)}$ based on (2):

$$(2^*) \beta^{(1)} = \rho^{(1)} + \zeta^{(1)}$$

→ Insert in (3)

$$\begin{aligned} (4) \|\beta^i\|^2 &= \sum_i^d (\rho^{(i)} + \zeta^{(i)})^2 \\ &= \sum_i (\rho^{(i)})^2 + 2\rho^{(i)}\zeta^{(i)} + (\zeta^{(i)})^2 \\ &= \|\rho\|_2^2 + 2\rho^T\zeta + \|\zeta\|_2^2 \end{aligned}$$

$2\rho^T\zeta = 0$ as $\rho \in \text{Row}(X)$ and $\zeta \in \text{Null}(X)$, which are orthogonal to each other

$$\Rightarrow \|\beta\|_2^2 = \|\rho\|_2^2 + \|\zeta\|_2^2$$

$$(4^*) \|\beta\|_2^2 = \|\rho\|_2^2 + \|\zeta\|_2^2 \geq \|\rho\|_2^2$$

With ζ being any vector in $\mathbb{R}^d$, we can say that the L2-norm of ρ is smaller or equal to the L2-norm of β .

(c) [5 points] **Coding question: minimum norm solution generalizes well**

For this sub-question, we still work with the setup of parts (a) and (b). We use the following datasets:

`src/implicitreg/ds1_train.csv, ds1_valid.csv`

Each file contains $d + 1$ columns. The first d columns in the i -th row represents $x^{(i)}$, and the last column represents $y^{(i)}$. In this sub-question, we use $d = 200$ and $n = 40$.

Using the formula in sub-question (b), **compute** the minimum norm solution using the training dataset. Then, **generate** three other different solutions with zero costs and different norms using the formula in sub-question (a). The starter code is in `src/implicitreg/linear.py`. **Plot** the validation error of these solutions (including the minimum norm solution) in a scatter plot. Use the norm of the solutions as x -axis, and the validation error as y -axis. For your convenience, the plotting function is provided as the method `generate_plot` in the starter code. Your plot is expected to demonstrate that the minimum norm solution generalizes well.

Answer:




Figure 8: Comparison of minimum norm solutions

- (d) [5 points] For this sub-question, we work with the setup of part (a) and (b). In this sub-question, you will prove that the gradient descent algorithm with *zero initialization* always converges to the minimum norm solution. Let $\beta^{(t)}$ be the parameters found by the GD algorithm at time step t . Recall that at step t , the gradient descent algorithm update the parameters in the following way

$$\beta^{(t)} = \beta^{(t-1)} - \eta \nabla J(\beta^{(t-1)}) = \beta^{(t-1)} - \frac{\eta}{n} X^\top (X\beta^{(t-1)} - \vec{y}). \quad (9)$$

As in sub-question (a), we also assume $XX^\top$ is an invertible matrix. **Prove** that if the GD algorithm with zero initialization converges to a solution $\hat{\beta}$ satisfying $J(\hat{\beta}) = 0$, then $\hat{\beta} = X^\top (XX^\top)^{-1} \vec{y} = \rho$, that is, $\hat{\beta}$ is the minimum norm solution.

Hint: As a first step, you can prove by induction that if we start with zero initialization, $\beta^{(t)}$ will always be a linear combination of $\{x^{(1)}, x^{(2)}, \dots, x^{(n)}\}$ for any $t \geq 0$. Then, for any $t \geq 0$, you can write $\beta^{(t)} = X^\top v^{(t)}$ for some $v^{(t)} \in \mathbb{R}^n$. As a second step, you can prove that if $\hat{\beta} = X^\top v^{(t)}$ for some $v^{(t)}$ and $J(\hat{\beta}) = 0$, then we have $\hat{\beta} = \rho$.

You don't necessarily have to follow the steps in this hint. But if you use the hint, you need to prove the statements in the hint.

Answer:

Loss Function: $J(\beta) = \frac{1}{n} \|X\beta - Y\|_2^2$

Gradient of Loss Function: $\frac{1}{n} X^\top (X\beta - Y)$

Minimum Norm Solution: $\hat{\beta} = X^\top (XX^\top)^{-1} Y$

Zero Loss solution (general): $\hat{\beta} = \rho + \zeta$ (For an invertible $XX^\top$)

1.

Assume $\beta^{(t)} = X^\top v^{(t)}$ for $v^{(t)} \in \mathbb{R}^n$

Proof for $t = 0, 1$:

$$\beta^{(0)} = 0 = X^\top \vec{0}$$

$$\beta^{(1)} = 0 - \eta \nabla J(\beta^{(0)}) = -\frac{\eta}{n} X^\top (X\beta^{(0)} - Y) = X^\top \overbrace{\frac{\eta}{n} (X^\top \beta^{(0)}) + Y}^{v^{(1)}}$$

Proof for $t = d + 1$:

$$\begin{aligned} \beta^{(d+1)} &= \beta^{(d)} - \eta \nabla J(\beta^{(d)}) \\ &= X^\top v^{(d)} - \eta \cdot \frac{1}{n} X^\top (X\beta^{(d)} - \vec{y}) \\ &= X^\top \cdot \underbrace{(v^{(d)} - \eta \cdot \frac{1}{n} (X\beta^{(d)} - \vec{y}))}_{v^{(d+1)}} \end{aligned}$$

2.

$$1) J(\hat{\beta}) = 0 \Leftrightarrow \frac{1}{2n} \|X\hat{\beta} - y\|_2^2 = 0 \Leftrightarrow X\hat{\beta} = y$$

$$2) \beta = X^\top v \text{ for some } v \in \mathbb{R}^n$$

Substituting 2) into 1):

$$X(X^\top v) = y \Rightarrow (XX^\top)v = y$$

$$v = (XX^T)^{-1}y$$

Insert into 2):

$$\beta = X^T(XX^T)^{-1}y = \rho$$

This proves that GD with zero initialization always ends up with the minimum norm solution, as it is constrained to the rowspace.

- (e) [3 points] In the following sub-questions, we consider a slightly more complicated model called quadratically parameterized model. A quadratically parameterized model has two sets of parameters $\theta, \phi \in \mathbb{R}^d$. Given a d -dimensional input $x \in \mathbb{R}^d$, the output of the model is

$$f_{\theta, \phi}(x) = \sum_{k=1}^d \theta_k^2 x_k - \sum_{k=1}^d \phi_k^2 x_k. \quad (10)$$

Note that $f_{\theta, \phi}(x)$ is linear in its input x , but non-linear in its parameters θ, ϕ . Thus, if the goal was to learn the function, one should simply just re-parameterize it with a linear model and use linear regression. However, here we insist on using the parameterization above in Eq. (10) in order to study the implicit regularization effect in models that are nonlinear in the parameters.

Notations: To simplify the equations, we define the following notations. For a vector $v \in \mathbb{R}^d$, let $v^{\odot 2}$ be its element-wise square (that is, $v^{\odot 2}$ is the vector $[v_1^2, v_2^2, \dots, v_d^2] \in \mathbb{R}^d$.) For two vectors $v, w \in \mathbb{R}^d$, let $v \odot w$ be their element-wise product (that is, $v \odot w$ is the vector $[v_1 w_1, v_2 w_2, \dots, v_d w_d] \in \mathbb{R}^d$.) Then our model can be written as

$$f_{\theta, \phi}(x) = x^\top (\theta^{\odot 2} - \phi^{\odot 2}). \quad (11)$$

Suppose we have a dataset $\{(x^{(i)}, y^{(i)}); i = 1, \dots, n\}$ where $x^{(i)} \in \mathbb{R}^d$ and $y^{(i)} \in \mathbb{R}$ for all $1 \leq i \leq n$, and

$$y^{(i)} = (x^{(i)})^\top ((\theta^*)^{\odot 2} - (\phi^*)^{\odot 2})$$

for some $\theta^*, \phi^* \in \mathbb{R}^d$. Similarly, we use $X \in \mathbb{R}^{n \times d}$ and $\vec{y} \in \mathbb{R}^n$ to denote the matrix/vector representing the inputs/labels respectively:

$$X = \begin{bmatrix} - & x^{(1)} & - \\ - & x^{(2)} & - \\ \vdots & \vdots & \vdots \\ - & x^{(n)} & - \end{bmatrix}, \quad \vec{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(n)} \end{bmatrix}.$$

Let $J(\theta, \phi) = \frac{1}{4n} \sum_{i=1}^n (f_{\theta, \phi}(x^{(i)}) - y^{(i)})^2$ be the cost function.

First, when $n < d$ and $XX^\top$ is invertible, **prove** that there exists infinitely many optimal solutions with zero cost.

Hint: Find a mapping between the parameter β in linear model and the parameter θ, ϕ in quadratically parameterized model. Then use the conclusion in sub-question (a).

Answer:

- i) $n < d$ (overparameterized regime)
- ii) $(XX^T)^{-1}$ exists

iii) Zero cost solution: $\beta = \rho + \zeta = X^T(XX^T)^{-1}Y + \zeta$

$$f_{\theta, \phi}(\theta) = X((\theta^*)^{\odot 2} - (\phi^*)^{\odot 2}) = Y$$

Let $\Theta = a\theta^* - (\theta^T p)$

$$f_{\theta, \phi}(\theta) = X\Theta = \vec{v}$$

→ This function $f_{\theta, \phi}$ has the same linear form from a) with $\vec{p} = Xp$, where $\beta = \theta$.

With the zero cost solution in iii) we can say:

$$\Theta = \phi + \zeta$$

Hence, same as in a), we can say that for Θ , which is dependent on θ, ϕ , there are infinitely many solutions for all $\zeta \in \text{Null}(X)$

(f) [10 points] **Coding question: implicit regularization of initialization**

We still work with the setup in part (e). For this sub-question, we use the following datasets:

`src/implicitreg/ds2_train.csv, ds2_valid.csv`

Each file contains $d + 1$ columns. The first d columns in the i -th row represents $x^{(i)}$, and the last column represents $y^{(i)}$. In this sub-question, we use $d = 200$ and $n = 40$.

First of all, the gradient of the loss has the following form:

$$\nabla_{\theta} J(\theta, \phi) = \frac{1}{n} \sum_{i=1}^n ((x^{(i)})^{\top} (\theta^{\odot 2} - \phi^{\odot 2}) - y^{(i)}) (\theta \odot x^{(i)}), \quad (12)$$

$$\nabla_{\phi} J(\theta, \phi) = -\frac{1}{n} \sum_{i=1}^n ((x^{(i)})^{\top} (\theta^{\odot 2} - \phi^{\odot 2}) - y^{(i)}) (\phi \odot x^{(i)}). \quad (13)$$

You don't need to prove these two equations. They can be verified directly using the chain rule.

Using the formula above, run gradient descent with initialization $\theta = \alpha \mathbf{1}, \phi = \alpha \mathbf{1}$ with $\alpha \in \{0.1, 0.03, 0.01\}$ (where $\mathbf{1} = [1, 1, \dots, 1] \in \mathbb{R}^d$ is the all-1's vector) and learning rate 0.08. We provide the starter code in `src/implicitreg/qp.py`. **Plot** the curve of training error and validation error with different α . Use the number of gradient steps as x -axis, and training/validation error as y -axis. Include your plot in the writeup and **answer** the following two questions based on your plot: which models can fit the training set? Which initialization achieves the best validation error?

Remark: Your plot is expected to demonstrate that the initialization plays an important role in the generalization performance—different initialization can lead to different global minimizers with different generalization performance. In other words, the initialization has an implicit regularization effect. **Answer:**




Figure 9: Validation loss with different initialization parameters

All of the models fit the training set pretty well but only the model that started with the lowest initial model parameters has the best generalization as it performs the best on the validation set.

(g) [6 points] **Coding question: implicit regularization of batch size**

We still work with the setup in part (e). For this sub-question, we use the same dataset and starter code as in sub-question (f). We will show that the noise in the training process also induces implicit regularization. In particular, the noise introduced by *stochastic* gradient descent in this case helps generalization. **Implement** the SGD algorithm, and **plot** the training and validation errors with batch size $\{1, 5, 40\}$, learning rate 0.08, and initialization $\alpha = 0.1$. Similarly, use the number of gradient steps as x -axis, and training/validation error as y -axis. For simplicity, the code for selecting a batch of examples is already provided in the starter code. **Compare** the results with those in sub-question (g) with the same initialization. Does SGD find a better solution?

Your plot is expected to show that the stochasticity in the training process is also an important factor in the generalization performance — in our setting, SGD finds a solution that generalizes better. In fact, a conjecture is that stochasticity in the optimization process (such as the noise introduced by a small batch size) helps the optimizer to find a solution that generalizes better. This conjecture can be proved in some simplified cases, such as the quadratically parameterized model in this sub-question (adapted from the paper HaoChen et al., 2020), and can be observed empirically in many other cases. **Answer:**




Figure 10: Validation loss with different batch sizes

SGD actually does find a better solution! For batchsize $n = 40$, we are equivalent to the initialization experiment setting with $\alpha = 0.1$, and the loss matches here. But the smaller our batches get, the better our validation loss performance. The noisier gradient updates seem to help us find even better minima, or at least solutions that generalize better.

5. [12 points] Double Descent on Linear Models

Background: In this question, you will empirically observe the sample-wise double descent phenomenon. That is, the validation losses of some learning algorithms or estimators do not monotonically decrease as we have more training examples, but instead have a curve with two U-shaped parts. The double descent phenomenon can be observed even for simple linear models. In this question, we consider the following setup. Let $\{(x^{(i)}, y^{(i)})\}_{i=1}^n$ be the training dataset. Let $X \in \mathbb{R}^{n \times d}$ be the matrix representing the inputs (i.e., the i -th row of X corresponds to $x^{(i)}$), and $\vec{y} \in \mathbb{R}^n$ the vector representing the labels (i.e., the i -th row of $\vec{y}$ corresponds to $y^{(i)}$):

$$X = \begin{bmatrix} - & x^{(1)} & - \\ - & x^{(2)} & - \\ \vdots & \vdots & \vdots \\ - & x^{(n)} & - \end{bmatrix}, \quad \vec{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(n)} \end{bmatrix}.$$

Similarly, we use $X_v \in \mathbb{R}^{m \times d}$, $\vec{y}_v \in \mathbb{R}^m$ to represent the validation dataset, where m is the size of the validation dataset. We assume that the data are generated with $d = 500$.

In this question, we consider *regularized* linear regression. For a regularization level $\lambda \geq 0$, define the regularized cost function

$$J_\lambda(\beta) = \frac{1}{2} \|X\beta - \vec{y}\|_2^2 + \frac{\lambda}{2} \|\beta\|_2^2,$$

and its minimizer $\hat{\beta}_\lambda = \arg \min_{\beta \in \mathbb{R}^d} J_\lambda(\beta)$.

- (a) [2 points] In this sub-question, we derive the closed-form solution of $\hat{\beta}_\lambda$. **Prove** that when $\lambda > 0$,

$$\hat{\beta}_\lambda = (X^\top X + \lambda I_{d \times d})^{-1} X^\top \vec{y} \quad (14)$$

(recall that $I_{d \times d} \in \mathbb{R}^{d \times d}$ is the identity matrix.)

Note: $\lambda = 0$ is a special case here. When $\lambda = 0$, $(X^\top X + \lambda I_{d \times d})$ could be singular. Therefore, there might be more than one solutions that minimize $J_0(\beta)$. In this case, we define $\hat{\beta}_0$ in the following way:

$$\hat{\beta}_0 = (X^\top X)^+ X^\top \vec{y}. \quad (15)$$

where $(X^\top X)^+$ denotes the Moore-Penrose pseudo-inverse of $X^\top X$. You don't need to prove the case when $\lambda = 0$, but this definition is useful in the following sub-questions.

Answer:

Global Minimum of $J_\lambda(\beta)$:

1) $\nabla J_\lambda(\beta) = X^\top (X\beta - Y) + \lambda\beta$

2)

$$\nabla J_\lambda(\beta) = 0 \Leftrightarrow 0 = X^\top (X\beta - Y) + \lambda\beta \quad (1)$$

$$0 = X^\top X\beta - X^\top Y + \lambda\beta \quad (2)$$

$$0 = X^\top X\beta - X^\top Y + \lambda I\beta \quad (3)$$

$$X^\top Y \frac{1}{\beta} = X^\top X + \lambda I \quad (4)$$

$$(X^\top Y)^{-1} \beta = (X^\top X + \lambda I)^{-1} \quad (5)$$

$$\beta = (X^T X + \lambda I)^{-1} X^T Y \quad (6)$$

Annotation for (3): For this formula to work, we need to multiply our identity value by λ to match the dimensions of $X^T X$. Thus this gives us the $\hat{\beta}_\lambda$ for $\lambda > 0$:

$$\hat{\beta}_\lambda = (X^T X + \lambda I_{d \times d})^{-1} X^T Y \quad (7)$$

(b) [5 points] **Coding question: the double descent phenomenon for unregularized models**

In this sub-question, you will empirically observe the double descent phenomenon. You are given 13 training datasets of sample sizes $n = 200, 250, \dots, 750$, and 800, and a validation dataset, located at

- `src/doubledescent/train200.csv`, `train250.csv`, etc.
- `src/doubledescent/validation.csv`

For each training dataset $(X, \vec{y})$, compute the corresponding $\hat{\beta}_0$, and evaluate the mean squared error (MSE) of $\hat{\beta}_0$ on the validation dataset. The MSE for your estimators $\hat{\beta}$ on a validation dataset $(X_v, \vec{y}_v)$ of size m is defined as:

$$\text{MSE}(\hat{\beta}) = \frac{1}{2m} \|X_v \hat{\beta} - \vec{y}_v\|_2^2.$$

Complete the `regression` method of `src/doubledescent/doubledescent.py` which takes in a training file and a validation file, and computes $\hat{\beta}_0$. You can use `numpy.linalg.pinv` to compute the pseudo-inverse.

In your writeup, include a line plot of the validation losses. The x-axis is the size of the training dataset (from 200 to 800); the y-axis is the MSE on the validation dataset. You should observe that the validation error increases and then decreases as we increase the sample size.

Note: When $n \approx d$, the test MSE could be very large. For better visualization, it is okay if the test MSE goes out of scope in the plot for some points.

Answer:




Figure 11: Double Descent phenomenon on different sample sizes

(c) [5 points] **Coding question: double descent phenomenon and the effect of regularization.**

In this sub-question, we will show that regularization mitigates the double descent phenomenon for linear regression. We will use the same datasets as specified in sub-question (b). Now consider using various regularization strengths. For $\lambda \in \{0, 1, 5, 10, 50, 250, 500, 1000\}$, you will compute the minimizer of $J_\lambda(\beta)$.

Complete the `ridge_regression` method of `src/doubledescent/doubledescent.py` which takes in a training file and a validation file, computes the $\hat{\beta}_\lambda$ that minimizes the training objective under different regularization strengths, and returns a list of validation errors (one for each choice of λ).

In your writeup, include a plot of the validation losses of these models. The x-axis is the size of the training dataset (from 200 to 800); the y-axis is the MSE on the validation dataset. Draw one line for each choice of λ connecting the validation errors across different training dataset sizes. Therefore, the plot should contain 8×13 points and 8 lines connecting them. You should observe that for some small λ 's, the validation error may increase and then decrease as we increase the sample size. However, double descent does not occur for a relatively large λ .

Remark: If you want to learn more about the double descent phenomenon and the effect of regularization, you can start with this paper Nakkiran, et al. 2020.

Answer:




Figure 12: Validation loss with different regularization parameters